



ATLAS CAMERA · SHOWCASE DEEP-DIVE 03 / 11 [← ALL WORKFLOWS](#)

THE CLASSIC DMP ANGLE, PART 1 — ANCHORED WALLS

Foreground buildings ground naturally; background bases are occluded and float. Geometry fixes what it can see.

FILE `atlas_dmp_angle_anchored_newyork_workflow.json` VERIFIED 2026-07-16

THE STORY

From a high window, foreground buildings meet the street in-frame — their geometry runs to $Y=0$ and they sit. Background buildings' bases hide behind nearer rooftops: the relief tears at the roofline and those meshes FLOAT in space, visibly detached the moment you orbit.

AtlasDeriveWalls with `ground_anchor=True` closes the gap wherever a facade VISIBLY meets the street: the footprint comes from pure rayground geometry (exact, regardless of depth-model error — the depth model is demoted to grouping pixels), and the wall extrudes FROM $Y=0$. `distance_modes=3` fits one wall per depth mode per facing direction — the street-grid skyline case, where a single median plane would collapse thirty facades into one slab.

The anchor deliberately REFUSES fully-occluded bases — its occlusion and contamination gates would rather leave a building floating than teleport it to a wrong distance. That residual is not a failure; it's exactly the gap Part 2 (the X-ray variant) exists to fill with predicted structure.

IN THE VIEWPORT — PROJECTED VS. THE GREY MESH




SOURCE PLATE

newyork_Birdseye.png · the street-contact problem — the single still photograph everything below is reconstructed from.




PROJECT ON

The photo reassembled on the recovered camera — texels assigned by ray, perfect from this view.




PROJECT OFF

The same geometry as a grey mesh, recovered view — what the depth-derived surfaces actually are.



GREY · ORBIT LEFT

Orbited off-axis: the reconstructed cone. Coverage ends where the camera never saw.



GREY · ORBIT RIGHT

The opposite angle — silhouette tears read as deliberate holes, not errors.

NODES ON STAGE

AtlasLearnedSolveFromImage

AtlasReferenceScaleSolve

AtlasDepthMap

AtlasDeriveWalls

AtlasDeriveReliefMesh

AtlasMergeGeometry

RETUNE ON YOUR PLATE

- **distance_modes (1–5).** More modes = more depth rows per facade direction. Raise on deep street grids; 1 restores the classic single-plane fit.
- **max_walls.** Dense skylines may need more than the default; each anchored row costs little.
- **Scale comes from counted storeys.** The anchor's rayground math needs the right metric ground plane — the counted-tenement reference (5×3.5 m) locks it at camera 63.7 m. Count levels on your own plate's reference building.
- **Best accuracy tip from the node docs:** inpaint cars/fences off the ground line before solving — visible ground contact is what the anchor trusts.

RUN-VERIFIED

RUN ✓ 13 anchored facades, every one planted at $Y \approx 0$, rising 22–54 m = 6–15 storeys at 3.5 m/level — countable against the photo, which independently validates the anchor geometry. Contamination teleports: zero.

REQUIREMENTS

Atlas + [neural] · no gates

