



ATLAS CAMERA · SHOWCASE DEEP-DIVE 01 / 11 [← ALL WORKFLOWS](#)

THE SOLVE LAB

Everything a camera solve carries, made visible — the teaching workflow.

FILE `atlas_solve_lab_coastal_alley_workflow.json` VERIFIED 2026-07-16

THE STORY

Before projecting a single pixel, you need to trust the solve. This workflow takes one plate with strong architectural lines — the classical vanishing-point solver's home turf — and lays out every diagnostic Atlas can produce: the VP fan-lines drawn over the photo, the horizon line with its confidence, ground and sky masks, an analytic ground depth map, and the decomposed camera numbers (focal, FOV, height, principal point).

It then walks the three-tier metric-scale ladder in one graph: a known-size reference object (a door, drawn as a bounding box) is tier 1; VLM-detected scale cues feed the same tier but only apply when you flip confirm ON; the assumed 1.6 m eye height is the flagged last resort. Finally, the solve is exported to JSON and RELOADED — the viewport renders the reloaded solve, proving the round-trip is lossless — alongside an artist-guided solve from explicit line groups and the legacy path-based node kept in a clearly-labeled corner.

IN THE VIEWPORT — PROJECTED VS. THE GREY MESH



SOURCE PLATE

coastal_alley_for_vlm.png · 1408×768 — the single still photograph everything below is reconstructed from.



PROJECT OFF

The same geometry as a grey mesh, recovered view — what the depth-derived surfaces actually are.



GREY · ORBIT LEFT

Orbited off-axis: the reconstructed cone. Coverage ends where the camera never saw.



SOLVE DIAGRAM



overlay: VP fan-lines, horizon and ground — the classical solve's evidence, drawn on the plate.

RETUNE ON YOUR PLATE

- ▶ **The reference bounding box is per-plate.** AtlasReferenceScaleSolve's bbox_x0/y0/x1/y1 are pixel coordinates for THIS door. New plate = redraw the box and pick the right reference_id (door / person / car).
- ▶ **The constraints JSON is per-plate too.** AtlasConstrainedSolve traces this alley's line groups; re-author the lines for a new photo or bypass the node.
- ▶ **VP solves need strong straight lines.** On organic or AI-generated plates the classical detector finds nothing — switch to AtlasLearnedSolveFromImage (GeoCalib).
- ▶ **VLM provider settings.** AtlasVLMScaleCues defaults to lmstudio :1234; point provider/base_url at your own server, or 'openai' for a cloud endpoint.

RUN-VERIFIED

RUN ✓ VP solve recovered an 85° hFOV; the door reference set camera height to 1.02 m; guided solve confidence 0.7. The VLM returned zero scale cues on this plate — and the confirm-gated apply correctly adopted nothing, which is the designed behavior.

REQUIREMENTS

Atlas + [neural] · VLM optional (fails soft) · no gates

